

Performance Modeling and Design of Computer Systems- Ch 1

Debobroto Das Robin

Kent State University

Spring 2020

drobin@kent.edu

Overview

Performance
Modeling and
Design of
Computer
Systems- Ch 1

Debobroto
Das Robin

Introduction

1 Introduction

Queueing Theory

Theory of Queues

Performance
Modeling and
Design of
Computer
Systems- Ch 1

Debobroto
Das Robin

Introduction

- The theory behind what happens when you have lots of jobs
- what makes queues appear and how to make them go away / Why queue size grows and shrinks?
- Queueing theory applies anywhere that queues come up
- Example
 - CPU uses a time-sharing scheduler to serve a queue of jobs waiting for CPU time
 - Router in a network serves a queue of packets waiting to be routed.
- Queueing theory is built on **stochastic modeling and analysis**
 - Model + analyze service demands of jobs and the interarrival times of jobs as random variables.

Goal of Queueing Theory

2 Goals

- Predicting the system performance. Ex.
 - predicting mean delay or delay variability in service
 - number of jobs that will be in queue
 - mean number of servers being utilized
- Developing design of improved system
- Example
 - Can we build a better system from 1 slow discs or one faster disc

Power of Queueing Theory

Design Example 1

- Consider a system of a single CPU that serves a queue of jobs in First-Come- First-Served (FCFS) order
- Assume any distribution for the arrival
 - λ = arrival rate = No of jobs arrives per second
 - μ = service rate = No of jobs served per second
 - **response time**= time diff. betn. job arrival time and until it completes service
 - $E[T]$ = mean response time $\sum(x.P(x))$.




Figure 1.2. A system with a single CPU that serves jobs in FCFS order.

- **Question:** What if λ is doubled?
 - How the $E[T]$ changes? increases or decreases?
 - Can you serve customer within old time length?
 - If decreases, can using powerful cpu solve the problem?
 - How much powerful CPU is necessary? ans: Less than double

Power of Queueing Theory

Design Example 2

Performance
Modeling and
Design of
Computer
Systems- Ch 1

Debobroto
Das Robin

Introduction

- Consider the closed system of figure




Figure 1.3. A closed batch system.

- Replace server 1 with a server that is twice as fast (the new server services jobs at an average rate of 2 jobs every 3 seconds).
 - Does this “improvement” affect the average response time in the system?
 - Does it affect the throughput?
 - Both cases improvement is no or negligible
- If the system is converted to open system? where arrival times are independent of service completions.
 - Absolutely possible

Power of Queueing Theory

Design Example 3

- one fast CPU of speed s , or n slow CPUs each of speed s/n (see Figure).
- Goal is to minimize mean response time.




Figure 1.5. Which is better for minimizing mean response time: many slow servers or one fast server?

- one fast CPU of speed s , or n slow CPUs each of speed s/n (see Figure 1.5). Your goal is to minimize mean response time.
 - Choice depends on **the variability of the job size**
 - **Question:** when job size variability is high? Answer: we prefer many slow servers because we do not want short jobs getting stuck behind long ones.
 - **Question:** Which system do you prefer when load is low? Answer: When load is low, not all servers will be utilized, so it seems better to go with one fast server.

Power of Queueing Theory

Design Example 3..cont..

Performance
Modeling and
Design of
Computer
Systems- Ch 1


Debobroto
Das Robin

Introduction

- If jobs are preemptible; Jobs can be stopped and restarted where they left off.
- **Question** do you prefer many slow machines as compared to a single fast machine? Answer: If your jobs are preemptible, you could always use a single fast machine to simulate the effect of n slow machines.
- Resources can vary. CPU, GPU, MEMORY etc.
- Complexity and variation grows

Power of Queueing Theory

Design Example 4



- Assume that all hosts are identical (homogeneous)
- all jobs only use a single resource.
- once jobs are assigned to a host, they are processed there in FCFS order and are non-preemptible.
- Which task assignment policies yields the lowest mean response time?
- Some options Random, Round-Robin, Size-Interval-Task-Assignment (SITA), Central- Queue:
- More possible. Answer depends on various parameters
 - If job size variability is low, then the LWL policy is best.
 - If job size variability is high, then it is important to keep short jobs from getting stuck behind long ones, so a SITA-like policy,

Power of Queueing Theory

Design Example 4

- A single server. Jobs arrive according to a Poisson process. Arbitrary Job size
- Question : Which Scheduling policy is best w.r.t. mean response time if **non- preemptive service orders**?
- Scheduling policies are: First-Come-First-Served (FCFS), Non-Preemptive Last-Come- First-Served (LCFS), etc
- Ans: All are same for **non-preemptive service orders**
- now if **Non Preemptive-LCFS policy (PLCFS)**-
Whenever a new arrival enters the system, it immediately preempts the job in service
- Mean response time depends on the variability of the job size distribution.
 - job size distribution is at least moderately variable, then PLCFS will be a huge improvement.

If the job size distribution is hardly variable (basically constant), then PLCFS policy will be up to a factor of 2 worse